

SwingScope

Administrator Guide

What runs by itself, what you must do, and what to do when something breaks. Total commitment: about two minutes a week.

29 July 2026

Research infrastructure. Not investment advice.

What you are actually administering

Five automated workflows and one irreplaceable file

The system runs itself. Your job is noticing when it stops, not operating it. Most weeks you check one thing and do nothing.

Admin — What Runs, and What You Do

AUTOMATIC — no action required

CI on every push · Mon 09:00 picks + Telegram · Weekdays 09:00 diary + Excel
Quarterly decay check · Monthly fundamentals snapshot

MONDAY about 2 minutes

- Did the Telegram message arrive?
- If not: Actions tab, open the run, find the red step

MONTHLY about 15 minutes

- Read the month-end report (arrives on the 1st)
- Confirm forward_log.csv grew by roughly 4 snapshots
- Check the F&O list and repo-rate table are not stale

QUARTERLY about 30 minutes

- Run the decay monitor
- Compare IC against 0.0311
- Upgrade dependencies one at a time, testing between each

AT 8 WEEKS the one that matters

- Read forward IC from the month-end report
- 40 to 60 percent retention = normal decay, still viable
- Near zero = the backtest did not hold. Stop.

NEVER DO THESE

Do not change parameters mid-experiment — it restarts the evidence clock

Do not merge past a red CI run

Do not put broker credentials in the repo or on Streamlit Cloud

The five workflows

File	Runs	Does	Writes
<code>ci.yml</code>	every push	syntax, imports, app wiring, 18 invariants, secret scan	—
<code>main.yml</code>	Mon 09:00 IST	picks, forward log, Telegram	<code>forward_log.csv</code>
<code>daily.yml</code>	weekdays 09:00	observation diary, Excel workbook	<code>daily_observations.csv</code>
<code>decay.yml</code>	quarterly	re-measures IC, alerts on decline	<code>decay_history.csv</code>
<code>fundamentals.yml</code>	monthly	point-in-time fundamentals snapshot	<code>fundamentals_archive.csv</code>

Known collision to fix. `main.yml` runs `30 3 * * 1` and `daily.yml` runs `30 3 * * 1-5` . Different strings, but both fire on Monday. You will get two Telegram messages and the two commits will race.

Change `daily.yml` to `30 3 * * 2-5` — Monday is already covered.

The file that matters

`forward_log.csv` is the only irreplaceable thing here.

Price data can be re-downloaded. Reports rebuild. Universes re-fetch. Forward evidence cannot be reconstructed, because it depends on having been recorded *before* the outcome was known. That is the entire point of it.

It is committed to git on every run, so history is recoverable via `git log -- forward_log.csv` .

Your tasks, in order

Every Monday — about two minutes

Did the Telegram message arrive?

That is the whole check. No message is the signal — and it looks identical to a quiet market, which is why the workflows carry a failure alert. If neither arrived, something failed before the alert could fire.

If nothing came

1. Actions tab — did the run start?
2. If red: open it, expand the failed step, read the error
3. If green but silent: look at `[10/10] Notify` for `telegram skipped`, which means the secrets did not reach the runner

Every month — about fifteen minutes

- **Read the month-end report**, which arrives on the 1st
- **Confirm the forward log grew** by roughly four snapshots. A stalled log is the quietest failure in the system
- **Check the hand-maintained tables** — the F&O list in `circuit.py` and the repo-rate history in `macro.py`. Both will silently persist wrong values if left

Every quarter — about thirty minutes

- Run the decay monitor, or let the scheduled job do it
- Compare IC against **0.0311**
- Upgrade dependencies **one at a time**, running `test_invariants.py 50` between each

Do not upgrade dependencies in bulk. Pinning was attempted once by hand and broke the deployment — the pins named versions well behind what the code actually ran on. If you want exact pins, generate them from a working install with `pip freeze`, never from memory.

At eight weeks — the decision

This is the only task that changes anything.

Forward IC vs 0.0311	Reading	Action
Above 0.019 (60%+)	Holding up well	Continue; consider small live size
0.012–0.019 (40–60%)	Normal decay	Continue paper trading
Below 0.012	Substantial decay	Do not size up
Zero or negative	The backtest did not hold	Stop. Investigate before resuming

Forty to sixty percent retention is *normal and healthy*, not a disappointment. Backtests are optimistic by construction because the signal was selected on the same data. A result that decays to nothing is the outcome this whole apparatus exists to catch — cheaply, on paper, rather than expensively with capital.

When something breaks

Admin — When Something Breaks

Symptom on the left, first thing to check on the right

No Telegram on Monday	Actions tab — did the run start at all?
Run shows red	Expand the failed step; the error is in its output
Green run, no message	Look at [10/10] Notify for 'telegram skipped'
Screener shows 0 passed	Read the filter attribution table — it names the binding filter
ABORT: no price data	yfinance rate limit. Wait 15 min, or the bhavcopy fallback covers it
CACHED FALLBACK in stage 1	NSE blocked the runner. Not urgent; the bundled list works
Health checks failed	This is the system working — it refused stale data
Decay alert fires	Do NOT change parameters. Check how many measurements exist first

The only irreplaceable file is `forward_log.csv`

Everything else regenerates. Forward evidence cannot — it had to be recorded before the outcome.

Disaster recovery

If the repository is lost:

1. The Streamlit deployment holds no state — nothing to recover there
2. Re-clone from any local copy or the distribution zip
3. Restore `forward_log.csv` from an Actions artifact (90-day retention) or a Telegram-delivered workbook
4. Re-add repository secrets — these are deliberately not in the repo
5. Re-enable Actions write permissions

If credentials leak:

1. **Telegram:** message @BotFather, `/revoke` , generate a new token

2. **Broker:** revoke the app in the developer portal immediately
3. Rotate *before* cleaning git history — the credential is already public, and removing it from history does not un-publish it

Recovery expectations

Scenario	Recovery
Missed weekly run	Manual trigger, ~5 min. One snapshot gap.
yfinance outage	Automatic bhavcopy fallback, no action
NSE blocks universe fetch	Automatic fallback to bundled list
Repository lost	~30 min from zip, plus secrets
Forward log lost	Not recoverable beyond the last commit

Things not to do

Do not change parameters mid-experiment

Horizon, universe size, bucket size, filters. Every change restarts the clock on clean forward evidence, and eight weeks is already the scarcest resource here.

Do not merge past a red CI run

The suite exists because manual verification repeatedly missed things during development — twice a patch silently failed to apply, and once two functions were deleted by a careless edit. CI catches all three classes in seconds.

Do not put broker credentials in the repo or on Streamlit Cloud

The repository is public for the free Streamlit tier. Bots scan public repos for credentials within minutes, and a committed secret lives in git history forever. Broker credentials belong on your own machine only.

Do not react to a single decay reading

Below three measurements the trend means nothing. Reacting to one reading is fitting noise, and the standard way systems get optimised into uselessness.

Do not intervene because a week was quiet

In a risk-off regime the bucket caps at three and often the honest answer is zero. Absence of trades is the design working.

Reference

Numbers to keep

Item	Value
Benchmark forward IC	0.0311 (20-year, 201 windows, $t = 2.74$)
Holding period	30 days
Universe	400 stocks, Nifty 500, liquidity-trimmed
Bucket size	10 (top 2.5%)
Expected net	+5.95% annualised after charges, slippage and tax
Experiment started	~28 July 2026

Manual commands

run.bat	launch the app locally
run_tests.bat 50	invariant suite
run_benchmark.bat	efficiency, precision, load, regression
run_validation_suite.bat	out-of-sample, Monte Carlo, risk metrics
run_pipeline.bat	the Monday job, on demand

Key files

File	Purpose
forward_log.csv	The evidence. Irreplaceable.
momentum.py	The signal and its validation record
pipeline.py	The ten-stage weekly job
governance.py	Model registry, provenance, audit trail
RUNBOOK.md	Fuller failure playbook

What is deliberately absent

No order placement. No execution layer. No FIX connectivity. The system computes stops and quantities and reports them; placing anything is manual.

That is not an oversight. With no forward evidence yet, the judgement step between a ranked list and an actual purchase is doing real work — and automating it would remove the thing currently protecting the account.

Research infrastructure. Not investment advice. The author is not a SEBI-registered investment adviser. Presenting this system's output as advice to third parties may require registration; take legal advice first.